



Deep Learning Basics

Lecture 10: Neural Language Models

Princeton University COS 495

Instructor: Yingyu Liang


Natural language Processing (NLP)

- The processing of the **human** languages by **computers**
- One of the oldest AI tasks
- One of the most important AI tasks
- One of the hottest AI tasks nowadays

Difficulty

- Difficulty 1: ambiguous, typically no formal description
- Example: “*We saw her duck.*”
- 1. We looked at a duck that belonged to her.
- 2. We looked at her quickly squat down to avoid something.
- 3. We use a saw to cut her duck.

Difficulty

- Difficulty 2: computers do not have human concepts
- Example: “She like *little animals*. For example, yesterday we saw her *duck*.”
- 1. We looked at a duck that belonged to her.
- 2. We looked at her quickly squat down to avoid something.
- 3. We use a saw to cut her duck.

Statistical language model

Probabilistic view

- Use probabilistic distribution to model the language
- Dates back to Shannon (information theory; bits in the message)

Statistical language model

- Language model: probability distribution over sequences of tokens
- Typically, tokens are words, and distribution is discrete
- Tokens can also be characters or even bytes
- Sentence: *“the quick brown fox jumps over the lazy dog”*

Tokens: x_1 x_2 x_3 x_4 x_5 x_6 x_7 x_8 x_9

Statistical language model

- For simplification, consider fixed length sequence of tokens (sentence)

$$(x_1, x_2, x_3, \dots, x_{\tau-1}, x_{\tau})$$

- Probabilistic model:

$$P [x_1, x_2, x_3, \dots, x_{\tau-1}, x_{\tau}]$$

N-gram model

n-gram model

- n -gram: sequence of n tokens
- n -gram model: define the conditional probability of the n -th token given the preceding $n - 1$ tokens

$$P[x_1, x_2, \dots, x_\tau] = P[x_1, \dots, x_{n-1}] \prod_{t=n}^{\tau} P[x_t | x_{t-n+1}, \dots, x_{t-1}]$$

n-gram model

- n -gram: sequence of n tokens
- n -gram model: define the conditional probability of the n -th token given the preceding $n - 1$ tokens

$$P[x_1, x_2, \dots, x_\tau] = P[x_1, \dots, x_{n-1}] \prod_{t=n}^{\tau} P[x_t | x_{t-n+1}, \dots, x_{t-1}]$$

Markovian assumptions



Typical n -gram model

- $n = 1$: unigram
- $n = 2$: bigram
- $n = 3$: trigram

Training n -gram model

- Straightforward counting: counting the co-occurrence of the grams

For all grams $(x_{t-n+1}, \dots, x_{t-1}, x_t)$

- 1. count and estimate $\hat{P}[x_{t-n+1}, \dots, x_{t-1}, x_t]$
- 2. count and estimate $\hat{P}[x_{t-n+1}, \dots, x_{t-1}]$
- 3. compute

$$\hat{P}[x_t | x_{t-n+1}, \dots, x_{t-1}] = \frac{\hat{P}[x_{t-n+1}, \dots, x_{t-1}, x_t]}{\hat{P}[x_{t-n+1}, \dots, x_{t-1}]}$$

A simple trigram example

- Sentence: “*the dog ran away*”

$$\hat{P}[\textit{the dog ran away}] = \hat{P}[\textit{the dog ran}] \hat{P}[\textit{away}|\textit{dog ran}]$$

$$\hat{P}[\textit{the dog ran away}] = \hat{P}[\textit{the dog ran}] \frac{\hat{P}[\textit{dog ran away}]}{\hat{P}[\textit{dog ran}]}$$

Drawback

- Sparsity issue: $\hat{P}[\dots]$ most likely to be 0
- Bad case: “*dog ran away*” never appear in the training corpus, so $\hat{P}[\textit{dog ran away}] = 0$
- Even worse: “*dog ran*” never appear in the training corpus, so $\hat{P}[\textit{dog ran}] = 0$

Rectify: smoothing

- Basic method: adding non-zero probability mass to zero entries
- Back-off methods: restore to lower order statistics
- Example: if $\hat{P}[\textit{away}|\textit{dog ran}]$ does not work, use $\hat{P}[\textit{away}|\textit{ran}]$ as replacement
- Mixture methods: use a linear combination of $\hat{P}[\textit{away}|\textit{ran}]$ and $\hat{P}[\textit{away}|\textit{dog ran}]$

Drawback

- High dimension: # of grams too large
- Vocabulary size: about $10^k = 2^{14}$
- #trigram: about 2^{42}

Rectify: clustering

- Class-based language models: cluster tokens into classes; replace each token with its class
- Significantly reduces the vocabulary size; also address sparsity issue
- Combinations of smoothing and clustering are also possible

Neural language model

Neural Language Models

- Language model designed for modeling natural language sequences by using a **distributed representation** of words
- Distributed representation: embed each word as a real vector (also called word embedding)
- Language model: functions that act on the vectors

Distributed vs Symbolic representation

- Symbolic representation: can be viewed as one-hot vector
- Token i in the vocabulary is represented as e_i

i -th entry




- Can be viewed as a special case of distributed representation

Distributed vs Symbolic representation

- Word embeddings: used for real value computation (instead of logic/grammar derivation, or discrete probabilistic model)
- Hope that **real value computation corresponds to semantics**
- Example: inner products correspond to token similarities
- One-hot vectors: every pair of words has inner product 0

Co-occurrence

- Firth's Hypothesis (1957): the meaning of a word is defined by “the company it keeps”




- Use the co-occurrence of the word as its vector:

$$v_w := \hat{P}[w, :]$$

Co-occurrence

- Firth's Hypothesis (1957): the meaning of a word is defined by “the company it keeps”



- Use the co-occurrence of the word as its vector:


$$v_w := \hat{P}[w, :]$$

Drawback

- High dimensionality: equal vocabulary size (~10k)
- can be even higher if context is used


Latent semantic analysis (LSA)

- LSA by [Deerwester et al., 1990](#): low rank approx. of co-occurrence



Variants

- low rank approx. of the **transformed** co-occurrence



$$\text{Or } \text{PMI}(w, w') = \ln \frac{\hat{P}[w, w']}{\hat{P}[w] \hat{P}[w']}$$


State-of-the-art word embeddings

Updated on April 2016

Word2vec

- Continuous-Bag-Of-Words

Figure from
*Efficient Estimation of Word
Representations in Vector Space*,
By Mikolov, Chen, Corrado, Dean



$$P[w_t | w_{t-2}, \dots, w_{t+2}] \propto \exp[v_{w_t} \cdot \text{mean}(v_{w_{t-2}}, \dots, v_{w_{t+2}})]$$

Linear structure for analogies


- Semantic: “man:woman::king:queen”

$$v_{man} - v_{woman} \approx v_{king} - v_{queen}$$

- Syntactic: “run:running::walk:walking”

$$v_{run} - v_{running} \approx v_{walk} - v_{walking}$$

Country and Capital Vectors Projected by PCA



GloVe: Global Vector

- Suppose the co-occurrence between word i and word j is X_{ij}
- The word vector for word i is w_i and \tilde{w}_i
- The GloVe objective function is

$$J = \sum_{i,j=1}^V f(X_{ij}) (w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log X_{ij})^2 ,$$

- Where b_i 's are bias terms, $f(x) = \min\{100, x^{3/4}\}$

Advertisement

Lots of mysterious things

What are the reasons behind

- The weird transformation on the co-occurrence?
- The model of word2vec?
- The objective of GloVe? The hyperparameters (weights, bias, etc)?

What are the connections between them? A unified framework?

Why do the word vector have linear structure for analogies?

Advertisement

- We proposed a generative model with theoretical analysis:
[RAND-WALK: A Latent Variable Model Approach to Word Embeddings](#)
- Next lecture by Tengyu Ma, presenting this work

Can't miss!